PLOS One

# Toward reliable false discovery rate control in classification problems under distribution shift

--Manuscript Draft--

| | |
|---|---|
| **Manuscript Number:** | PONE-D-26-04885 |
| **Article Type:** | Research Article |
| **Full Title:** | Toward reliable false discovery rate control in classification problems under distribution shift |
| **Short Title:** | Toward reliable false discovery rate control in classification problems under distribution shift |
| **Corresponding Author:** | Attila Kertesz-Farkas<br>National Research University Higher School of Economics: Nacional'nyj issledovatel'skij universitet Vyssaa skola ekonomiki<br>Moscow, Moscow RUSSIAN FEDERATION |
| **Keywords:** | Machine Learning, Error control, |
| **Abstract:** | Background: Domain shifts or batch effects can significantly degrade the performance of machine learning-based classifiers in software production without any warning. This can particularly be problematic in medical applications which involve human lives. The shift changes the shape of the prediction score distribution of the test samples compared to that of the training samples; hence, decision boundary re-calibration would be required to conform to the expected performance standards.<br>Method and Findings: We present a simple, yet robust heuristic method to calculate valid p-values (uniformly distributed) for raw prediction scores of test samples under data shift in general classification problems. Our method is based on an adjustment of the test prediction scores so that the test null distribution (approximated distribution of the prediction scores of negative test samples) aligns to the training null distribution (actual distribution of prediction scores of the negative training samples). We termed it the Test Null Adjustment (TNA) method. Then, empirical p-values of the adjusted test prediction scores can be calculated with respect to the training null distribution to be used in any statistical applications which require test samples accompanied by valid p-values. We demonstrate the usability of our TNA method in FDR control with Benjamini-Hochberg (BH) protocol.<br>Conclusion: The main advantage of the TNA method is that its application does not require any domain specific knowledge. It works in binary and multi-label classification scenarios, it is fully data-driven, and it operates in the 1-dimensional space of the prediction scores ensuring stable numerical calculations. We have performed four experimental tests on the biomedical image analysis domain to demonstrate that the TNA method is able to mitigate the bias in FDR induced by data shift in real-life applications. |
| **Order of Authors:** | Attila Kertesz-Farkas |
| | Andrey Borevskiy |
| **Opposed Reviewers:** | |
| **Additional Information:** | |
| **Question** | **Response** |
| **Financial Disclosure**<br><br>Enter a financial disclosure statement that describes the sources of funding for the work included in this submission. Review the submission guidelines for detailed requirements. View published research articles from *PLOS ONE* for specific | The author(s) received no specific funding for this work. |

*Powered by Editorial Manager® and ProduXion Manager® from Aries Systems Corporation*

| | |
|---|---|
| examples.<br><br>This statement is required for submission and **will appear in the published article** if the submission is accepted. Please make sure it is accurate.<br><br>**Funded studies**<br>Enter a statement with the following details:<br>• Initials of the authors who received each award<br>• Grant numbers awarded to each author<br>• The full name of each funder<br>• URL of each funder website<br>• Did the sponsors or funders play any role in the study design, data collection and analysis, decision to publish, or preparation of the manuscript?<br><br>Did you receive funding for this work? | |
| **Competing Interests**<br><br>Use the instructions below to enter a competing interest statement for this submission. On behalf of all authors, disclose any competing interests that could be perceived to bias this work—acknowledging all financial support and any other relevant financial or non-financial competing interests.<br><br>This statement is required for submission and **will appear in the published article** if the submission is accepted. Please make sure it is accurate and that any funding sources listed in your Funding Information later in the submission form are also declared in your Financial Disclosure statement.<br><br>View published research articles from *PLOS ONE* for specific examples. | The authors have declared that no competing interests exist. |

| | |
|---|---|
| **NO authors have competing interests**<br><br>Enter: *The authors have declared that no competing interests exist.*<br><br>**Authors with competing interests**<br><br>Enter competing interest details beginning with this statement:<br><br>*I have read the journal's policy and the authors of this manuscript have the following competing interests: [insert competing interests here]*<br><br>* typeset | |
| **Ethics Statement**<br><br>Enter an ethics statement for this submission. This statement is required if the study involved:<br><br>• Human participants<br>• Human specimens or tissue<br>• Vertebrate animals or cephalopods<br>• Vertebrate embryos or tissues<br>• Field research<br><br>Write "N/A" if the submission does not require an ethics statement.<br><br>General guidance is provided below. Consult the submission guidelines for detailed instructions. **Make sure that all information entered here is included in the Methods section of the manuscript.** | N/A |

| | |
|---|---|
| **Format for specific study types**<br><br>**Human Subject Research (involving human participants and/or tissue)**<br>• Give the name of the institutional review board or ethics committee that approved the study<br>• Include the approval number and/or a statement indicating approval of this research<br>• Indicate the form of consent obtained (written/oral) or the reason that consent was not obtained (e.g. the data were analyzed anonymously)<br><br>**Animal Research (involving vertebrate animals, embryos or tissues)**<br>• Provide the name of the Institutional Animal Care and Use Committee (IACUC) or other relevant ethics board that reviewed the study protocol, and indicate whether they approved this research or granted a formal waiver of ethical approval<br>• Include an approval number if one was obtained<br>• If the study involved *non-human primates*, add *additional details* about animal welfare and steps taken to ameliorate suffering<br>• If anesthesia, euthanasia, or any kind of animal sacrifice is part of the study, include briefly which substances and/or methods were applied<br><br>**Field Research**<br><br>Include the following details if this study involves the collection of plant, animal, or other materials from a natural setting:<br>• Field permit number<br>• Name of the institution or relevant body that granted permission | |
| **Data Availability**<br><br>Authors are required to make all data underlying the findings described fully available, without restriction, and from the time of publication. PLOS allows rare exceptions to address legal and ethical concerns. See the PLOS Data Policy and FAQ for detailed information. | Yes - all data are fully available without restriction |

| | |
|---|---|
| A Data Availability Statement describing where the data can be found is required at submission. Your answers to this question constitute the Data Availability Statement and **will be published in the article**, if accepted.<br><br>**Important:** Stating 'data available on request from the author' is not sufficient. If your data are only available upon request, select 'No' for the first question and explain your exceptional situation in the text box.<br><br>Do the authors confirm that all data underlying the findings described in their manuscript are fully available without restriction? | |
| **Describe where the data may be found in full sentences. If you are copying our sample text, replace any instances of XXX with the appropriate details.**<br><br>• If the data are **held or will be held in a public repository**, include URLs, accession numbers or DOIs. If this information will only be available after acceptance, indicate this by ticking the box below. For example: *All XXX files are available from the XXX database (accession number(s) XXX, XXX.*).<br>• If the data are all contained **within the manuscript and/or Supporting Information files**, enter the following: *All relevant data are within the manuscript and its Supporting Information files.*<br>• If neither of these applies but you are able to provide **details of access elsewhere**, with or without limitations, please do so. For example:<br><br>*Data cannot be shared publicly because of [XXX]. Data are available from the XXX Institutional Data Access / Ethics Committee (contact via XXX) for researchers who meet the criteria for access to confidential data.*<br><br>*The data underlying the results presented in the study are available from (include the name of the third party* | This study used existing data from public repository. URLs to data sources are provided in the main text in Table 1. |

| | |
|---|---|
| *and contact information or URL).*<br>• This text is appropriate if the data are owned by a third party and authors do not have permission to share the data.<br><br>* typeset | |
| Additional data availability information: | |

Managing Editor
*Plos One*

12-Jan-2026

Dear Editor:

We are submitting our manuscript entitled “Toward reliable false discovery rate control in classification problems under distribution shift” for consideration for publication in *Plos One* as ***original research article***.

*We note that we are affiliated to a Russian University, our institution is not under sanctions and we are not involved in any research related to the Russian military, etc. We will pay APC from a European Union (Hungarian) bank card.*

We present our study on controlling false discovery rates (FDR) under data distribution shift in classification problems in the biomedical domain. In this study, we draw attention to a very important but often neglected problem; namely, the performance of deep-learning based classifiers is assumed to be the same in application. However, it often can be much worse than thought due to data distribution shift (a.k.a. covariate shift), data label distribution shift, batch effects caused by some confounding factors, and overfitting. This performance drop of a machine learning system often remains undetected, resulting in risk for people in health care and clinical applications.

In this manuscript, we present a simple, fast and robust method to adjust the test prediction scores so that the latent test null distribution will be more similar to the train null distribution. This will result in more accurate FDR control.

Our method operates in the 1-dimensional prediction score space. Therefore, it is not hindered by high-dimensional space, data sparsity (i.e. curse-of-dimensionality), and correlated features. In addition, our method does not rely on data clustering or graphical causal models.

We tested our method with four public benchmark datasets commonly used in biomedical image analysis: PatchCamelyon (PCam), Chest x-ray (CheXpert), TissueNet, and Breast Cancer Semantic Segmentation (BCSS). With our method, we obtained much more accurate FDR control under data distribution shifts in the test data compared to a recent algorithm, called Learn-then-test, developed by Michael Jordan’s lab from Stanford University.

**A declaration of any potential competing interests:**

None.

**Confirmation that all authors have approved the manuscript for submission:**

This manuscript has been approved by both authors.

**Confirmation that the content of the manuscript has not been published:**

This manuscript has never been submitted to any other journals. We have not interacted with PLOS regarding this manuscript before. The manuscript will be properly formatted for the journal if it is needed for the reviewing stage or once it is accepted. We did not use AI chat bots in manuscript preparation or in this study in general.

**Suggested Academic Editors:**

1. Musa Aydin, Turkey, orcid.org/0000-0002-5825-2230
2. Hong-Seng Gan, Malaysia, orcid.org/0000-0003-3777-3640
3. Zeyneb Kurt, UK, orcid.org/0000-0003-3186-8091
4. Xiaohui Zhang, USA, orcid.org/0000-0002-7990-1594
5. Anybody with a background in machine learning classification, error control, image processing would suffice.

**Suggested reviewers:**

1. Nikita Moshkov, PhD, Institute of Biochemistry, Biological Research Center. Szeged, Hungary, Deep Learning specialist in biological image analysis, cell segmentation, etc. nikita.moshkov@brc.hu
2. Peter Horvath, PI at AI for Health, Helmholtz Munich, Germany; PI at Biological Research Center, Szeged, Hungary horvath.peter@brc.hu
3. Pavel Sulimov, PhD, ZHAW School of Engineering, Switzerland, pavel.sulimov@zhaw.ch
4. Yang Lu, PhD, Assistant Professor, University of Waterloo, Canada, yanglu@uwaterloo.ca
5. Luca Oneto, full Professor at University of Genoa, luca.oneto@unige.it

**Opposed reviewers:**

None.

We hope that this manuscript provides a useful and interesting study for the readers of Plos One.

Best regards,

Attila Kertesz-Farkas,
Professor
Head of the Laboratory on AI for Computational Biology, Faculty of Computer Science
HSE University, Moscow, Russia

Manuscript
Click here to access/download;Manuscript;exact-p-value_v3_clean.pdf

# Toward reliable false discovery rate control in classification problems under distribution shift

Andrey Borevskiy[1], Attila Kertesz-Farkas[1,*],

**1** Laboratory on AI for Computational Biology, Faculty of Computer Science, HSE University, 11 Pokrovsky Bvld, Moscow, 109028, Russian Federation

∗ kfattila@yandex.com

## Abstract

**Background:** Domain shifts or batch effects can significantly degrade the performance of machine learning-based classifiers in software production without any warning. This can particularly be problematic in medical applications which involve human lives. The shift changes the shape of the prediction score distribution of the test samples compared to that of the training samples; hence, decision boundary re-calibration would be required to conform to the expected performance standards.

**Method and Findings:** We present a simple, yet robust heuristic method to calculate valid p-values (uniformly distributed) for raw prediction scores of test samples under data shift in general classification problems. Our method is based on an adjustment of the test prediction scores so that the test null distribution (approximated distribution of the prediction scores of negative test samples) aligns to the training null distribution (actual distribution of prediction scores of the negative training samples). We termed it the Test Null Adjustment (TNA) method. Then, empirical p-values of the adjusted test prediction scores can be calculated with respect to the training null distribution to be used in any statistical applications which require test samples accompanied by valid p-values. We demonstrate the usability of our TNA method in FDR control with Benjamini-Hochberg (BH) protocol.

**Conclusion:** The main advantage of the TNA method is that its application does not require any domain specific knowledge. It works in binary and multi-label classification scenarios, it is fully data-driven, and it operates in the 1-dimensional space of the prediction scores ensuring stable numerical calculations. We have performed four experimental tests on the biomedical image analysis domain to demonstrate that the TNA method is able to mitigate the bias in FDR induced by data shift in real-life applications.

## Introduction

Modern AI applications not only need to increase their discriminative power to reduce the overall number of incorrect predictions, but also need to simultaneously minimize the quantified risks of false positive (FP) and negative (FN) error by an optimal decision threshold calibration while also being able to handle (or detect) data distribution shift [1, 2]. False positive (type I) and false negative (type II) errors can have significantly different costs or impacts on people and society [3]. For instance, a false positive HIV, COVID-19 or cancer diagnosis can cause emotional trauma and the person may undergo invasive treatments with undesirable complications and side effects [4–7] until a more elaborate clinical test establishes the correct diagnosis.

However, a failure to diagnose an infectious disease may lead to a false sense of safety [8], and may put not only their own but also other people's lives in danger [9]. An undetected cancer may either delay or hinder treatment [10], and misdiagnosed human papillomavirus (HPV) may progress to cancer precursor lesions [11, 12].

False discovery rate (FDR) control procedures are plausible techniques to keep the FP rate among the total positive classifications at a certain, user-defined level [13], when the risk of FP prediction is significantly higher than the cost of FN predictions. FDR control became popular with the emergence of high-throughput and shotgun experiments which measured large numbers of distinct variables simultaneously per sample in a cost efficient way. For instance, automatic patch clamp systems employ a deep-learning-based image processing module to automatically detect target cells (e.g. nerve cells in living brain). Then it calibrates the movement of a micropipette to approach the cell for future experiments [14]. Cell morphological profiling in high-throughput imaging assays for studying compound libraries and human diseases [15] relies on automatic cell segmentation and detection often in 3D images [16]. The FDR control here could curb the error in selecting wrong cells for subsequent experiments and feature quantification. Microarrays measure the expression levels of thousands of genes simultaneously and the FDR control allows the selection of promising genes for followup studies [17–19]. Tandem mass spectrometry can be used to quantify and identify thousands of proteins in a complex biological/chemical sample and FDR control makes it possible to choose the most trustworthy protein identifications [20].

The False Selection Rate (FSR) control procedure can restrain the FP and FN errors simultaneously [21]. This approach determines two thresholds: one to make definitive positive predictions, and another one to make definitive negative predictions. Any samples not passing either threshold remain in *indecision*. Therefore, the misclassification error (accuracy) among definitive predictions can be bounded, while a subset of samples without predictions are returned for manual investigation by a human expert or for further data collection [22]. Typically, the FDR and FSR controls, and many other statistical methods heavily rely on valid, i.e. uniformly distributed p-values, while typical deep learning models produce raw discriminative scores or logits for test samples.

In this paper, we are interested in providing valid p-values for a set of discriminative scores $T_s = \{t_1, t_2, \ldots, t_m\}$, where each score $t_i$ is produced by a machine learning classifier for a test instance $t_i$ without knowing its true label. The scores $t_i$s are from a mixture of unknown null and non-null hypotheses; i.e. $t_i$ is generated under one of the two possible distributions: (1) null distribution: $T_i \sim T_0$, and (2) non-null distribution: $T_i \sim T_1$; in a mixture model: $T_i \sim T = \pi_0 T_0 + (1 - \pi_0)T_1$, where $\pi_0$ is the prior probability of that a score comes from the null distribution. The null distribution is associated with the negative class, the non-null distribution is associated with the positive class. The null and the non-null distributions as well as the assignments are unknown due to the lack of information on the true class labels. Implicitly or explicitly, but knowing or estimating the null distribution is essential to obtain accurate, valid p-values with respect to it. In certain practical applications it is possible to label some of the data as negative with high confidence. For instance, decoy peptides [20] in tandem mass spectrometry and internal negative control proteins [23] in proteomic profiling serve this purpose. However, this is not the case with general machine learning applications, and we propose to approximate the test null distribution $T_0$ with the empirical training null $X_0 = \{x_1, x_2, \ldots, x_n\}$, where all $x_i$s prediction scores belong to the negative training class; that is, $T_0 \approx X_0$.

Therefore, an *empirical p-value* (EPV) of a given discriminative score $t$ could be defined as the fraction of the scores of the negative training data that are greater than or equal to $t$, formally $p_t = (1 + \sum_{x_i \in X_0} \mathbb{1}\{x_i \geq t\})/(n + 1)$ [eq. ]. This empirical

p-value is also referred to as conformal p-value [24] or as rank among negative controls (RANC) [23]. The EPVs of the negative test samples are super uniform under the null if the data distribution of the test and training samples are identically and independently distributed. This property is called *exchangeability*. In this case, the Benjamini-Hochberg (BH) procedure controls the FDR at $\pi_0\alpha$ level, where $\alpha$ is a user-defined confidence level. The EPVs have been successfully used in recent methods for outlier detection [24, 25] and in classification [22, 26] under the assumption of exchangeability.

Unfortunately, the validity of p-values and EPVs can be compromised by various problems including data distribution shift (a.k.a. covariate shift), data label distribution shift, batch effects [27] caused by some confounding factors, and overfitting. As a consequence multiple testing procedures that rely on valid p-values may fail to control the relevant statistical errors. In natural sciences, this shift can be attributed including, but not limited, to alterations in laboratory or environmental conditions, change in the reagents or atmospheric ozone levels, as well as to personnel preferences and to variation of technical sources [28]. Various high-throughput instruments produced by different manufacturers may also produce slightly differently distributed data about the same experiments. Quantitative morphological features of cells from multi-channel fluorescence microscopy images may depend on lamp intensities, filter patterns, and dyes [29]. Tandem mass spectrometry data are also affected by differences in experimental protocols and data acquisition conditions, reagent batches, or changes in instrumentation [30, 31]. These impacts may lead to incorrect conclusions about the experiments and to improper clinical drug therapies, etc. Similarly in economics and industry, machine learning-based predictions may also become inaccurate after economic shocks, technological advancement, geopolitical turbulence, and pandemic [32]. For instance, energy consumption prediction needs to be recalibrated after energy price change in order to maintain carbon emission reasonably low [33]; credit risk assessment needs to be adapted to monetary policy tightening, increased government expenditure, and income distribution shift caused by inflation [34–36].

Statistically speaking, the data shift violates the exchangeability assumption resulting in different test and training null hypotheses ($T_0 \not\approx X_0$) that further results in invalid p-values and conservatively or liberally (anti-conservatively) biased FDR control. Recent methods, which deal with data shift, aim to (1) detect data shift [37], (2) identify features in tabular data, which cause data shift [38], (3) detect and correct label shift [39], (4) provide explainability about data shift [40, 41], and (5) aim to adapt to data shifts [42–44], respectively. Most of these methods operate in the feature space and rely on data clustering or graphical causal models. They are also challenged by high-dimensional space, data sparsity (i.e. curse-of-dimensionality) [45], and correlated features.

In this paper, we propose a method to calculate EPVs for raw test scores under data or class distribution shift in general classification problems. These EPVs can be used in any statistical applications which rely on valid p-values. We demonstrate the usability of our method in FDR control with BH procedure. Our method, first, approximates the unknown test null distribution $T_0$ relying on the known training null and non-null distributions. Then, our method adjusts the test prediction scores so that the adjusted test null distribution $\hat{T_0}$ aligns well to the training null distribution; that is, $\hat{T_0} \approx X_0$. Hence, we termed our method Test Null Adjustment (TNA). Fig 1 illustrates this idea. The empirical p-values of the adjusted test scores can be calculated with respect to the training null distribution $X_0$ with eq [1]. The test null adjusted empirical p-values (TNA EPVs) can then be used by the user in any statistical methods that require valid p-values, including $\pi_0$ estimation, or FDR control with any procedure. Valid p-values should be uniformly distributed over the range of $[0, 1]$. We propose to visually monitor

**Fig 1. Sketch of TNA.** (A) The plot illustrates the distribution shift of the discriminative scores of the test (blue) instances from the training (red) instances. The shaded areas illustrate the observed training and unobserved test null distributions, respectively. The black arrow indicates that the test scores should be shifted to leftward. The TNA method aims to approximate the unknown test null distribution (blue shaded area) based on the known true training null distribution (red shaded area). (B) TNA aims to perform a test score adjustment so that the test null aligns to training null distributions better.

the uniformity of the p-values with Q-Q plots, in which the p-values are plotted against their ordered normalized positions. Any deviation from the diagonal line of the Q-Q plot would indicate a data shift.

The TNA method operates solely at the space of prediction score, which is typically one dimensional for each class. Therefore, TNA is not hindered by the curse-of-dimensionality [45]. It can be employed with any machine learning model, given that the classification is based on discriminative scores. This includes deep neural networks, black box models, support vector machines [46], with any kind of data such as tabular, text, audio, image, graph data. The TNA method is data driven as it does not rely on any class of analytical distributions. The only assumption we have for the TNA method is that the score distributions "somewhat resemble" the ones in Fig 2A or in Fig 10 that are typical.

# Materials and methods

## Preliminaries

Let us suppose that we are given domain $D$ such as images, text, speech, graphs, time series, tabular data, or the mix of these; furthermore, let $X \subseteq D$ and $T \subseteq D$ denote the training and the test data samples, respectively. Let $Y_t \in \{0, 1\}$ denote the true class label of sample $t \in D$. Let $f : D \to R$ be an appropriately trained binary discriminator (classifier), such as a deep neural network, to predict positive instances, where $R$ denotes the real numbers. Without loss of generality, we assume that higher score $f(t)$ indicates a stronger membership to the positive class. The empirical training null distribution is constructed as $X_0 = \{f(t) : Y_t = 0,\ t \in X\}$ from the scores of negative instances, the empirical training target distribution is constructed as $X_1 = \{f(t) : Y_t = 1,\ t \in X\}$ from the scores of positive instances, and the training score set is $X_s = \{f(t) : t \in X\}$. The true test null $T_0$, and true test target $T_1$ are unobserved, but the observed test score set is $T_s = \{f(t) : t \in T\}$. We say that we trust (or accept) a sample $t$ predicted as positive if and only if $f(t) > s_t$ for a given decision threshold $s_t$ and we make indecision for samples if $f(t) < s_t$. The false discovery proportion (FDP) in a set of samples $K = \{t \in D\}$ at a score threshold $s_t$ is calculated as

$$FDP(s_t, K) = \frac{1 + |\{t \in K : Y_t = 0,\ f(t) \geq s_t\}|}{|\{t \in K : f(t) \geq s_t\}|},$$

where $|.|$ denotes set cardinality. We reserve FDR for the statistical expectation of FDP. FDR control methods aim to calibrate the threshold $s_t$ so that the error rate among trusted predictions is at most level of $\alpha$.

The empirical p-values (EPVs) of the test samples are calculated using eq. with respect to $X_0$. The p-value calculation can be time consuming for large $X_0$ sets, in this case, a binned, empirical cumulative distribution function of $X_0$ could be used to

calculate the p-values accordingly in an accelerated way. Then, the FDR of a set $T$ can be controlled using the $BH$ protocol and the EPVs at a predetermined $\alpha$ level. For the sake of simplicity, we replace the $\alpha$ with $\alpha/\hat{\pi_0}$ inside BH algorithm, where $\hat{\pi_0}$ is an estimation of $\pi_0$ introduced by Storey [47] in the following way:
$\hat{\pi_0} = (1 + \sum_{i=1}^{m} \mathbb{1}\{p_i \geq \lambda_S\})/(m(1 - \lambda_S))$ for $\lambda_S \in (0, 1)$. The q-value of a test sample $t \in T$ is defined as $q_t = \min_{z \leq f(t)} FDP(z, T)$; that is the minimum $\alpha$ level when the sample $t$ would become a trusted positive prediction. The smaller the q-value of a test sample $t$ is, the more trusted positive prediction $t$ is. Note that, the q-value of a sample depends not only on its prediction score, but also on the other samples in the set.

### Learn-then-test (LTT).

We acquired the LTT code from its authors' GitHub repository. We specified the following hyperparameters as following: 100 uniformly distributed queried levels of FDR control are iteratively being processed, starting from the 1e-3, with the most appropriate threshold $\lambda_L$ being selected among the 300 potential threshold values for each iteration. The remainder of the methodology is entirely dependent on the published pipeline, including the comprehensive calculation of the p-values for each threshold during each alpha level test, with a further Bonferroni procedure employed to select the most appropriate border. The final value of the test false positive rate is equal to the number of samples with higher scores than the selected $\lambda_L$.

### Example 1.

We used the well-known MNIST dataset for demonstration [48] (see Table 1). It contains 60K training and 10K test images of 28x28 grey-scaled pixels in 10 classes. There is no data or label shift between the training and test data. We trained a convolutional neural network to classify the digit '1' as target (positive) against all the other digits (negative) (binary classification for the sake of simplicity). Our ConvNet model contained two convolutional layers consisting of 8 kernels with size of $3 \times 3$ each, and it was trained with an Adam optimizer with a learning rate of 0.01 for three epochs. Our trained model achieved a 99% accuracy on the test set.

The relevant results are presented in Fig 2. Fig 2A shows the training and test null and target distributions. Fig 2B shows a Q-Q plot of the test EPVs against the theoretical uniform distribution; that is the scatter plot of the EPV against its normalized rank in our case. The dots line up along the diagonal line which shows that the EPVs calculated are uniformly distributed and, consequently, valid. This also indicates that there is no shift between the training and test distributions. Fig 2C shows the number of accepted (trusted) predictions as positives at various q-values ($\alpha$ level of FDR). The ground truth (black solid line) was calculated using the true test labels. Any FDR controlling method (a) yielding more trusted predictions than the ground truth is implied to be liberally biased, and (b) yielding less trusted predictions is conservatively biased. The BH procedure with EPVs (red dashed line) yields closely the same number of trusted predictions at various FDR levels of $\alpha$ (calculated without using the true test labels). This repeatedly indicates that the EPVs are uniform and the BH procedure is an unbiased method. The LTT method also accurately controls the FDR at any $\alpha$ levels (blue dotted line). Finally, Fig 2D shows that the deviation of the estimated FDPs from the true FDPs is small.

### TNA protocol.

When the empirical training null distribution $X_0$ is different from the unobserved, true test null distribution $T_0$, the EPVs will not be uniformly distributed and the BH

**Fig 2. FDR control in MNIST with EPVs.** (A) The discriminative score distributions. (B) The Q-Q plot of the EPVs of the test samples against the theoretical uniform distribution. (C) The number of the accepted (trusted) positive predictions at various q-values calculated with true labels (black solid line), BH with EPVs (red dashed line), and LTT (blue dotted line). (D) Deviation between the true and the estimated FDR produced by BH with EPV (red line) and LTT (blue dots), respectively.

procedure results in liberal or conservative bias in FDR estimation. In the following steps, we describe two versions of TNA methods that both adjusts the prediction scores so that the approximated test null distribution $\hat{T}_0$ aligns to $X_0$. The first variation, called TNA- (TNA minus), relies on the assumption that the shapes of the training null ($X_0$) and test null ($T_0$) distributions are "similar", but there is no assumption on the form of the non-null distributions. The second variation, called TNA+ (TNA plus), relies on that the shapes of the training non-null ($X_1$) and test non-null ($T_1$) distributions are "similar", but there is no assumption on the analytical form of the null distributions. By "similar" distributions we mean that they are close in a statistical sense, such as having a small Wasserstein distance of a Jensen-Shannon divergence; intuitively, their histograms can be aligned without major reshaping. TNA's performance is robust to minor distributional shifts in practice, but it degrades gracefully as this distance increases. The proportion of the nulls may freely vary between the training and test data in the case of both TNA methods, there is no assumption on this. The two approaches differ only at the 3rd step of the following algorithm.

1. Get the density histograms of $X_0$, $X_1$, and $T_s$ distributions with identical bin borders and denote them as $\overline{X}_0$, $\overline{X}_1$, and $\overline{T}_s$, respectively. The $\overline{X}[i]$ denotes the value of bin $i$. We used 100 bins in our experiments.
2. Let $\hat{\pi}_{0_T}$ be an estimation of the proportion of the negative test predictions calculated as follows: $\hat{\pi}_{0_T} = N/(N+P)$, where N and P denotes the total number of the negative and positive predictions in the given test data, respectively, as predicted by the classifier $f$. We note that this $\hat{\pi}_{0_T}$ estimation can be done with raw discriminative scores and it does not require valid p-values.

3a. TNA-: We create a new density histogram $\hat{T}_0$ that will be an approximation of the test null $T_0$. The value of the histogram bin $i$ is calculated by $\hat{T}_0[i] = \overline{T}_s[i] \cdot c[i]$, where $c[i] = X_0[i]/X_s[i] \cdot \pi_{0_X}/\hat{\pi}_{0_T}$ and $\pi_{0_X}$ is the actual proportion of the null samples in the training data calculated with using the true class labels.

3b. TNA+: This approach relies on the following general formulation: $A = (\pi_0) \cdot A_0 + (1-\pi_0) \cdot A_1$; therefore, the null distribution can be expressed as $A_0 = (A - (1-\pi_0) \cdot A_1)/\pi_0$. We create a new density histogram $\hat{T}_0$ that will be an approximation of the test null $T_0$. The value of the histogram bin $i$ is calculated by $\hat{T}_0[i] = (\overline{T}_s[i] - (1-\hat{\pi}_{0_T})\hat{T}_1[i])/\hat{\pi}_{0_T}$. The unknown $\hat{T}_1$ is approximated by $T_s$ scaled with the proportion of the target calculated from the training data; that is:

$$\hat{T}_1[i] = \overline{T_s}[i] \cdot \frac{\overline{X_1}[i]}{\hat{\pi}_{0_T}\overline{X_0}[i] + (1-\hat{\pi}_{0_T})\overline{X_1}[i]}.$$

Note that the empirical *training* null and target distributions $(X_0, X_1)$ are weighted with *test* class proportion $\hat{\pi}_{0_T}$. The outcome of division by zero is set to zero.

4. Let $\mu_{X_0}$, $\sigma_{X_0}$, $\mu_{\hat{T_0}}$, and $\sigma_{\hat{T_0}}$ are the mean and std of the $X_0$ and the $\hat{T}_0$ distributions, respectively.
5. Adjust the test scores $\hat{t_s} = (t_s - \mu_{\hat{T_0}})/\sigma_{\hat{T_0}} \cdot \sigma_{\hat{X_0}} + \mu_{X_0}$.

Finally, calculate the TNA EPVs for each adjusted $\hat{t_s}$ test score with respect to the empirical training null ($X_0$) with using Eq. [1]. The TNA EPVs then can be used in any statistical applications that rely on valid p-values. In our experiments in the following sections, we demonstrate the usability of the TNA EPVs in FDR control with BH procedure. With TNA EPVs, we re-estimate $\hat{\pi_0}$ for the test instances with the Storey method with $\lambda_S = 0.8$, and run BH procedure to control the FDR at a desired level.

We note that the actual number of bins or the size of bin-width in step 1 may depend on the range and/or the number of the discriminative scores. Essentially, there should be enough data in each bins to ensure accurate estimations. This might require the the analyst to manually adjust the bins in practice.

We also note that we tested a one-dimensional optimal transportation (OT) to adjust training null to test null. Unfortunately, this approach is very sensitive to the test null approximation especially at the decision boundary. In our experiment, OT severely overfitted, while TNA remained more robust to poor tail approximation of the null distributions. OT did not provide any good results (data not shown) in our benchmark.

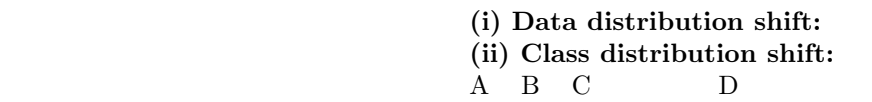


**Fig 3. FDR control under data (top row) and label (bottom row) shift.** (A) Q-Q plot of the EPVs (red dots), the TNA- EPVs (green dots) and the TNA+ EPVs (purple dots) against the theoretical uniform distribution (normalized rank). (B) The number of accepted classifications as a function of the Q-values over a critical range (0-0.1) obtained with (i) ground truth (black line), (ii) BH with EPV (red), (iii) BH with TNA- EPV (green), (iv) BH with TNA- EPV (purple), and (v) LTT (blue). (C) Same as (B) but over the entire q-value range. (D) Deviation of the estimated FDP from the true FDP obtained with (i) BH with EPV (red line), (ii) BH with TNA- EPV (green line), (iii) BH with TNA+ EPV (purple line), and (iv) LTT (blue line).

**Example 2.**

To illustrate the case of data distribution shift, we reduced the pixel intensity by 10% of the test images in the MNIST dataset (but training images remained the same). The accuracy of the CNN classifier (the same as trained in Example ) remained at 99 %. The evaluation plots are shown in Fig 3(i). The Q-Q plot (Fig. 3(i)A) shows that the test EPVs (red) became biased due to data distribution shift; however, both TNA methods can correct the EPVs yielding uniformly distributed p-values (green, purple). The number of accepted predictions ((Fig. 3(i)BC) obtained with BH with EPVs (red) differs from the ground truth substantially (black line), it became liberally biased. The LTT yielded similar amount of accepted predictions (blue) compared to the ground truth in the critical range of q-values ((Fig. 3(i)B); however, LTT also becomes liberally biased at lower ranges of q-values ((Fig. 3(i)C). The number of accepted predictions obtained with BH with both TNA EPVs (green, purple) remained close to the ground truth over both the critical and lower ranges of q-values implying accurate FDR control under data distribution shift. The actual deviation of the estimated FDP from the true FDP is shown in Fig 3(i)D for each method. The TNA+ is slightly less accurate than TNA- in this test scenario.

**Example 3.**

The FDR estimation is sensitive to changes in class proportions. To illustrate the case of class distribution shift, we resampled the test data so that the number of true negative

and positive instances became equal. This approach was used in other articles, e.g. [49]. The accuracy of the CNN classifier (the same as trained in Example ) remained at 99 %. The evaluation plots are shown in Fig 3(ii). The Q-Q plot (Fig. 3(ii)A) shows that the EPVs (red) remained unbiased and so did TNA EPVs (green, purple). This is in fact expected, because the shift in class distributions should not affect the actual null distributions; Therefore, the results of BH-based methods rely on the $\hat{\pi_0}$ estimation. BH with EPV and TNA EPV remained accurate and produced results close to the ground truth (red, green, and purple lines in Fig 3(ii)BCD) indicating successful $\hat{\pi_0}$ estimation. The BH procedure is robust to changes in class proportions because it can involve the (re-)estimation of the negative data proportion ($\hat{\pi_0}$) that can be done with, e.g. Storey method. The LTT is also sensitive to changes in class proportions because it relies on the class proportions that are calculated with the training (validation) datasets, and it assumes that class proportions remain unchanged in test phase. Here, LTT resulted in conservative FDR control even at the critical range of q-values (blue in Fig. 3(ii)BC). The bias in FDR produced by LTT (blue) is also visible in Fig 3(ii)D.

# Results

The datasets analyzed in this study are listed in Table 1.

**Table 1.** Data Availability Statement containing dataset names, repositories, and accession details.

| Dataset | Full Name | Accession / URL |
|---|---|---|
| MNIST | MNIST Database of Handwritten Digits | DOI: 10.24432/C53K8Q; `http://yann.lecun.com/exdb/mnist/` |
| BCSS | Breast Cancer Semantic Segmentation Dataset | DOI: 10.1093/bioinformatics/btz083; `https://github.com/PathologyDataScience/BCSS` |
| CheXpert | CheXpert Chest X-ray Dataset | `https://aimi.stanford.edu/datasets/chexpert-chest-x-rays` |
| PatchCamelyon | PatchCamelyon (PCam) Benchmark Dataset | DOI: 10.5281/zenodo.2546921 |
| TissueNet | TissueNet Cell Segmentation Dataset | DOI: 10.1038/s41587-021-01094-0; `https://datasets.deepcell.org/` |

## PCam: tumor classification

The PatchCamelyon (PCam) benchmark dataset is a collection of histopathologic scans of lymph node sections [50] derived from the Camelyon 16 challenge [51] (see Table 1). PCam consists of 327,680 colored images of size 96x96 each. An arbitrarily chosen sample data is shown in Fig 4. The challenge is to identify the presence (positive) or absence (negative) of any histopathology in one image. The dataset creators of PCam emphasize that all the positive and negative instances in the training, validation, test sets are equally distributed, resulting in a class balance of 1:1. We used the pre-trained resnet34-pcam model, which is a popular, deep residual network from the TIA (Tissue Image Analysis) toolbox [52] with an overall f1-score of 0.889, as stated by the authors. We evaluated the uniformity of the p-values of the test samples and the performance of the FDR controlling methods. The results are shown in the first row of Fig 5. The Q-Q plots (Fig. 5-(i)-A) of the test EPVs (red dots) indicate that the test p-values are

slightly biased, possibly indicating a slight overfitting and/or distribution shift between the training and test datasets. However, the Q-Q plots of the test EPVs adjusted with TNA- (green dots) and TNA+ (purple dots) show a more uniformly distributed p-values. The plots on Fig. 5-(i)-BCD indicate that all four FDR control methods are mostly accurate with this dataset especially at critical $\alpha$ levels (0-0.1), despite the small bias in the EPVs. The TNA methods (both, TNA- and TNA+) manage to reduce the bias from EPVs and yield more accurate FDR control with BH (green and purple lines), than the BH method with the original EPVs (red line). However, LTT outperformed both BH-based methods in this benchmark.

**Fig 4. Instances from PCam dataset.**

In our opinion, the actual class proportions in real-life applications can be significantly different from 1:1, making the FDR control with methods like LTT inaccurate in practical applications. We manually resampled the test data so that the positive and negative class balance became 31%:69%, following the original class balance of the Camelyon 16 dataset. The evaluation plots are shown in the second row of Fig 5. The Q-Q plot (Fig. 5-(ii)-A) reveals that the EPVs (red dots) become slightly biased; however, our TNA methods manage to reduce this bias from the EPVs (green dots). The LTT method became liberal in FDR control, the BH protocol with standard EPVs resulted in conservative FDR control; whereas, the BH with the TNA EPVs remained accurate, especially at critical range of q-values (0-0.1) (Fig. 5-(ii)-C).

**(i) Standard PCam dataset:**
**(ii) PCam with class distribution shift:**
**(iii) Standard CheXpert dataset:**
**(iv) CheXpert dataset with class distribution shift:**
A B C D

**Fig 5. Performance evaluation of FDR controlling methods.** See the caption of Fig 3 for plot descriptions.

**Fig 6. Instances of CheXpert dataset.**

## CheXpert: chest x-ray dataset

The CheXpert dataset (see Table 1) contains a total of 224,316 radiographs of both frontal and lateral views of the thorax from ∼65K patients in order to facilitate the automated interpretation of chest X-rays for multi-label classification tasks (i.e. multiple labels can be predicted for each instances) with a total of 14 labels [53]. An arbitrarily chosen sample data is shown in Fig 6. We used the deep neural network model, called DenseNet121, from the `Torchxrayvision` package [54, 55], specifically developed for this CheXpert dataset. For the sake of simplicity, we turned the CheXpert dataset into a binary classification task by treating all instances of pleural effusion as positive and all the others as negative. We selected this label because this was the most prevalent disease in the dataset. The proportion of the positive class became 27 % and 23 % in the training and test datasets, respectively. The DenseNet121 model achieved accuracies of 89 % and 84 % on the training and test sets, respectively.

We evaluated the four FDR controlling methods and plotted the results on Fig 5-(iii). Interestingly, the most confident prediction of the DenseNet121 model in the test set is wrong. This can be seen by the flat region of the ground truth evaluation (black

line) at low q-values around 0-0.05 (Fig, 5-(iii)-C). Fortunately, all four FDR controlling methods under investigations managed to capture this and remain close to the ground truth (Fig, 5-(iii)-B-D). The standard EPVs (red dots) tend to be off at the critical range (p-value $< 0.1$), but this does not affect the FDR controlling results in this benchmark, perhaps because of the confident incorrect predictions. Overall, we conclude that all four methods provide accurate FDR control. In our opinion, this indicates the absence of any database distribution shits. These plots also show that our TNA- and TNA+ algorithms do not corrupt the EPVs in the absence of data shift.

The actual frequency of pleural effusion may be different among all chest x-ray radiographs performed. It may vary over locations, age groups, in-patients and out-patients [56, 57]. In the United States, approximately 275 million conventional radiology procedures are performed annually [58], but only around 1.5 million of US citizens develop pleural effusion [57]. Pleural effusion develops in about 20-40% of in-patient with pneumonia [59]. Whatever the real abundance of the pleural effusion is among chest x-rays performed, it can be significantly different from the 27 % of the CheXpert dataset. Taking this into account, we re-sampled the test set in the CheXpert dataset to achieve a lower positive class ratio of 13%. The results of the re-evaluation are shown in Fig 5-(iv). The LTT and the TNA- EPVs are both slightly biased. The BH with the standard EVPs and with TNA+ EPVs achieve the most accurate FDR control over the entire range of $\alpha$ levels (0-1) in general. The LTT method is very sensitive to class distribution shift demonstrated by this case, wheree it resulted in a liberally biased FDR control.

## TissueNet: cells segmentation

**Fig 7. Instance of TissueNet dataset.** Two-channel original images are being processed by the model, in turn producing masks & flows to correctly segment cells

TissueNet dataset contains 1.3 million microscopic images of cells obtained from six various platforms and nine organs including both histologically healthy and diseased tissues [60](see Table 1). Each images has size of 512x512 pixels, was manually segmented at the pixel level and paired whole-cell and nuclear annotations. An arbitrarily chosen sample data is shown in Fig 7. Our binary classification task was to predict whether a pixel is inside a cell (positive) or not (negative). The overall positive-negative ratio was 59%:41% in the training data and in contrast to 66%:34% in the test data. However, the class ratio varies over the images considerably. We used the Cellpose pre-trained deep neural network model [61] and fine-tuned it for five epochs. It achieved accuracy of 95% and 94% on training and test data, respectively. Our evaluation is reported in Fig 9-(i). The plots reveal that the BH protocol with standard EPVs becomes liberally biased (Fig. 9-(i)-B-C). The latter possibly originates from the data distribution shift or from overfitting. Both LTT and the BH with TNA EPV remain accurate.

The situation changes when we evaluate the segmentation problem with a single image. We arbitrarily chose one test image (ID: 007-008-009) and run the Cellpose model to classify each pixel. The positive and negative class balance was: 33%:67%. The Q-Q plot (Fig. 9-(ii)-A) of the test EPVs (red dots) indicates that the scores of negative samples from training and test datasets are different. The TNA EPVs on the Q-Q plot shows that the TNA methods have managed to adjust the p-values of the test samples (green and purple dots) closer to the uniform distribution. The plots on Fig. 9-(ii)-B-D indicate that LTT has turned out to be liberally biased, the BH method with standard EPVs strongly conservative. The BH with TNA- EPVs (green) and with TNA+ EPVs

(purple) remain accurate in the critical region of q-values (0-0.4). Thus, the TNA methods have succeeded in correcting the p-values at the critical range at least.

## BCSS: Breast cancer semantic segmentation dataset

**Fig 8. An instance from BCSS dataset.**

The Breast Cancer Semantic Segmentation (BCSS) dataset [62] contains around 20K image segmentation objects annotated through Digital Slide Archive, with a particular focus on tissue regions derived from breast cancer imagery obtained from The Cancer Genome Atlas (TCGA)(see Table 1). The image pixels are labeled by: 'Stroma' (class id: 0), 'Tumor' (class id: 1) and 'Others' (class id: 2). An arbitrarily chosen sample data is shown in Fig 8. The corresponding class ratios in the training and test sets are 56%:38%:6% and 33%:59%:8%, respectively. We used the deep learning model called `'fcn resnet50 unet-bcss'` from the TIA (Tissue Image Analysis) toolbox [52] to classify the pixels to one of the three classes. The model achieved an accuracy of 84 % and 82 % on the training and test datasets, respectively.

Since this dataset involves a multi-class classification problem in which the FDR control is not straightforward due to the competition nature of the class prediction; thus, we slightly modified the FDR control procedure. We calculated the p-values, the $\hat{\pi_0}$ estimations with Storey method with respect to a single class (deemed as positive); and repeat it for each class via one-vs-rest fashion. Thus, for each pixels, we got three p-values, each corresponding to one of the three classes. Then we concatenated all of these p-values and calculated the q-values. We note the FDR control in this case allows the prediction of a single pixel into multiple classes. This is essentially a multi-label classification scenario.

The evaluation of the FDR controlling methods is shown in Fig 9-(iii). At first glance, the results suggest that LTT manages to control FDR well but this datasets is challenging for BH with standard or TNA EPVs. The Q-Q plots of the EPVs and the TNA EPVs are far from being uniform (Fig. 9-(iii)-A). Overall, the BH control with standard and TNA EPVs is conservative (Fig. 9-(iii)-B), but it is slightly liberal over the critical range of q-values (0-0.1) (Fig. 9-(iii)-C). BH protocol seems to be more accurate with standard rather than with TNA EPVs. The FDR estimation error with BH is significant (Fig. 9-(iii)-D).

Next, we further analyzed the results by each class separately. We plotted the training and test positive and negative discriminative score distributions class-wise provided by the deep learning model (Fig. 10 first row). The distribution plots show that the test scores are differently distributed from the training scores in all the classes. We also plotted the distribution of the adjusted scores of the negative test samples, adjusted by our TNA methods (green and purple lines in Fig. 10) in order to see how well they are aligned to the training null distribution. In the case of the Stroma class, the mode of adjusted test scores (green and purple lines) move a bit closer to the one of the true training null distribution (X0, filled orange bars). However, their tails on the right hand side remain off. In the other two classes, the adjusted test scores align visibly

**(i) Standard TissueNet dataset:**
**(ii) TissueNet with single image:**
**(iii) Standard BCSS dataset:**
A B C D

**Fig 9. Performance evaluation of FDR controlling methods.** See the caption of Fig 3 for plot descriptions.

better to the training null distribution (X0). The corresponding FDR curves are shown in the second row of Fig 10. Both the TNA- and the TNA+ methods improves the results in the case of the Stroma class. This is mainly due to the fact that the Storey method provides better $\hat{\pi_0}$ estimation using the TNA EPVs than using the standard EPVs. In the case of the other two classes, TNA- and TNA+ make minor effects on the FDR control compared to the standard EPV. The LTT method is overly liberal with the Stroma class, and it is overly conservative with the Tumor class. We suspect that LTT appears to be accurate by chance on the whole dataset because the two types of bias cancel each other out in the Stroma and the Tumor classes. These two classes account for the the 94 % of the whole data.

Stroma Tumor Others

**Fig 10. Classwise score distributions and FDR control in the BCSS dataset.**

## Conclusion

In this study, we examined the effects of data shift on False Discovery Rate (FDR) control. In our opinion, the BH procedure with standard EPVs provides a very accurate FDR control in the absence of any data shift. Unfortunately, data shift is rather common in practice, and we demonstrated that the standard BH and the LTT methods are quite sensitive even to small data or class distribution shift. In this paper, we presented a simple, but robust heuristic method, called Test Null Adjustment (TNA), in order to mitigate the effect of the data shift on the FDR control. Our method operates in the space of prediction scores, thereby it is not affected by the curse-of-dimensionality problem, which often hinders methods that aim to mitigate data shift in the input feature space. The TNA method consists of two main steps. First, it approximates the null distribution among the scores of the test instances via certain statistics from the scores of the training instances. Second, it recalibrates the test scores so that the test and training null distributions would move closer to each other. TNA is fully data-driven, it does not rely on any theoretical assumptions or forms of the data. We have demonstrated in our experimental test with four biomedical datasets that TNA can successfully mitigate the bias caused by data shift; while, leaving the distributions intact in the absence of data shift. We note that, the handling of data shifts in practice is challenging, and unfortunately our method does not provide any theoretical guarantee that the FDR control with TNA EPVs becomes accurate, or conservatively or liberally biased without any additional assumptions; hence TNA is heuristic in its general form.

With this study, we also aim to emphasize and popularize the importance of accurate error control in any predictions in the biomedical domain. We have shown that the deviations in the simple Q-Q plots of the calculated vs. the theoretical p-values can be used to detect data shifts between the training and test data. Moreover, we also hope that the FDR control and plots with q-values will become a standard metric in performance evaluation for prediction systems in the biomedical domain.

## References


1. Feng J, Phillips RV, Malenica I, Bishara A, Hubbard AE, Celi LA, et al. Clinical artificial intelligence quality improvement: towards continual monitoring and updating of AI algorithms in healthcare. npj Digital Medicine. 2022;5(1):66.
2. Al-Antari MA. Artificial Intelligence for Medical Diagnostics—Existing and Future AI Technology!. MDPI; 2023.

3. Wynants L, Van Smeden M, McLernon DJ, Timmerman D, Steyerberg EW, Van Calster B, et al. Three myths about risk thresholds for prediction models. BMC medicine. 2019;17:1-7.

4. Newman-Toker DE, Wang Z, Zhu Y, Nassery N, Tehrani ASS, Schaffer AC, et al. Rate of diagnostic errors and serious misdiagnosis-related harms for major vascular events, infections, and cancers: toward a national incidence estimate using the "Big Three". Diagnosis. 2021;8(1):67-84.

5. Salz T, Richman AR, Brewer NT. Meta-analyses of the effect of false-positive mammograms on generic and specific psychosocial outcomes. Psycho-Oncology. 2010;19(10):1026-34.

6. Tosteson ANA, Fryback DG, Hammond CS, Hanna LG, Grove MR, Brown M, et al. Consequences of False-Positive Screening Mammograms. JAMA Internal Medicine. 2014 06;174(6):954-61. Available from: https://doi.org/10.1001/jamainternmed.2014.981. arXiv:https://jamanetwork.com/journals/jamainternalmedicine/articlepdf/1861037/ioi140 doi:10.1001/jamainternmed.2014.981.

7. Xu Y, Nguyen D, Mohamed A, Carcel C, Li Q, Kutlubaev MA, et al. Frequency of a false positive diagnosis of epilepsy: a systematic review of observational studies. Seizure. 2016;41:167-74.

8. Mouliou DS, Gourgoulianis KI. False-positive and false-negative COVID-19 cases: respiratory prevention and management strategies, vaccination, and further perspectives. Expert review of respiratory medicine. 2021;15(8):993-1002.

9. Woloshin S, Patel N, Kesselheim AS. False negative tests for SARS-CoV-2 infection—challenges and implications. New England Journal of Medicine. 2020;383(6):e38.

10. Bradley SH, Nicholson BD, Funston G. Interpreting negative test results when assessing cancer risk in general practice. British Journal of General Practice. 2021;71(708):298-9.

11. Macios A, Nowakowski A. False Negative Results in Cervical Cancer Screening—Risks, Reasons and Implications for Clinical Practice and Public Health. Diagnostics. 2022;12(6):1508.

12. Pinsky PF. Principles of cancer screening. Surgical Clinics. 2015;95(5):953-66.

13. Benjamini Y, Hochberg Y. Controlling the false discovery rate: a practical and powerful approach to multiple testing. J Roy Statist Soc Ser B. 1995;57(1):289-300. Available from: http://links.jstor.org/sici?sici=0035-9246(1995)57:1<289:CTFDRA>2.0.CO;2-E&origin=MSN.

14. Koos K, Oláh G, Balassa T, Mihut N, Rózsa M, Ozsvár A, et al. Automatic deep learning-driven label-free image-guided patch clamp system. Nature communications. 2021;12(1):1-11.

15. Moshkov N, Bornholdt M, Benoit S, Smith M, McQuin C, Goodman A, et al. Learning representations for image-based profiling of perturbations. Nature Communications. 2024;15(1):1594.

16. Falk T, Mai D, Bensch R, Çiçek Ö, Abdulkadir A, Marrakchi Y, et al. U-Net: deep learning for cell counting, detection, and morphometry. Nature methods. 2019;16(1):67-70.

17. Storey JD. A direct approach to false discovery rates. Journal of the Royal Statistical Society Series B: Statistical Methodology. 2002;64(3):479-98.

18. Storey JD, Tibshirani R. Statistical significance for genomewide studies. Proceedings of the National Academy of Sciences. 2003;100(16):9440-5.

19. Barber RF, Candès EJ. Controlling the false discovery rate via knockoffs. The Annals of statistics. 2015:2055-85.

20. Elias JE, Gygi SP. Target-decoy search strategy for increased confidence in large-scale protein identifications by mass spectrometry. Nature methods. 2007;4(3):207.

21. Zhao G, Su Z. Controlling FSR in Selective Classification. arXiv preprint arXiv:231103811. 2023.

22. Rava B, Sun W, James GM, Tong X. A burden shared is a burden halved: A fairness-adjusted approach to classification. arXiv preprint arXiv:211005720. 2021.

23. Gao Z, Zhao Q. Simultaneous hypothesis testing using internal negative controls with an application to proteomics. arXiv preprint arXiv:230301552. 2023.

24. Bates S, Candès E, Lei L, Romano Y, Sesia M. Testing for outliers with conformal p-values. The Annals of Statistics. 2023;51(1):149-78.

25. Marandon A, Lei L, Mary D, Roquain E. Adaptive novelty detection with false discovery rate guarantee; 2023. arXiv:2208.06685.

26. Angelopoulos AN, Bates S, Candès EJ, Jordan MI, Lei L. Learn then test: Calibrating predictive algorithms to achieve risk control. arXiv preprint arXiv:211001052. 2021.

27. Quionero-Candela J, Sugiyama M, Schwaighofer A, Lawrence ND. Dataset Shift in Machine Learning. The MIT Press; 2009.

28. Leek JT, Scharpf RB, Bravo HC, Simcha D, Langmead B, Johnson WE, et al. Tackling the widespread and critical impact of batch effects in high-throughput data. Nature Reviews Genetics. 2010;11(10):733-9.

29. Bray MA, Singh S, Han H, Davis CT, Borgeson B, Hartland C, et al. Cell Painting, a high-content image-based assay for morphological profiling using multiplexed fluorescent dyes. Nature protocols. 2016;11(9):1757-74.

30. Phua SX, Lim KP, Goh WWB. Perspectives for better batch effect correction in mass-spectrometry-based proteomics. Computational and Structural Biotechnology Journal. 2022.

31. Čuklina J, Lee CH, Williams EG, Sajic T, Collins BC, Rodríguez Martínez M, et al. Diagnostics and correction of batch effects in large-scale proteomic studies: A tutorial. Molecular systems biology. 2021;17(8):e10240.

32. Ramey VA. Macroeconomic shocks and their propagation. Handbook of macroeconomics. 2016;2:71-162.

33. Clement T, Nguyen HTT, Kemmerzell N, Abdelaal M, Stjelja D. Coping with Data Distribution Shifts: XAI-Based Adaptive Learning with SHAP Clustering for Energy Consumption Prediction. In: Australasian Joint Conference on Artificial Intelligence. Springer; 2023. p. 147-59.

34. Kritzman M, Page S, Turkington D. Regime shifts: Implications for dynamic strategies (corrected). Financial analysts journal. 2012;68(3):22-39.

35. Guo Y, Hu C, Yang Y. Predict the Future from the Past? On the Temporal Data Distribution Shift in Financial Sentiment Classifications. arXiv preprint arXiv:231012620. 2023.

36. Zhang M. Adaptation Based Approaches to Distribution Shift Problems. EECS Department, University of California, Berkeley; 2021. Available from: `http://www2.eecs.berkeley.edu/Pubs/TechRpts/2021/EECS-2021-262.html`.

37. Dasu T, Krishnan S, Lin D, Venkatasubramanian S, Yi K. Change (detection) you can believe in: Finding distributional shifts in data streams. In: Advances in Intelligent Data Analysis VIII: 8th International Symposium on Intelligent Data Analysis, IDA 2009, Lyon, France, August 31-September 2, 2009. Proceedings 8. Springer; 2009. p. 21-34.

38. Kulinski S, Bagchi S, Inouye DI. Feature shift detection: Localizing which features have shifted via conditional distribution tests. Advances in neural information processing systems. 2020;33:19523-33.

39. Lipton Z, Wang YX, Smola A. Detecting and correcting for label shift with black box predictors. In: International conference on machine learning. PMLR; 2018. p. 3122-30.

40. Budhathoki K, Janzing D, Bloebaum P, Ng H. Why did the distribution change? In: International Conference on Artificial Intelligence and Statistics. PMLR; 2021. p. 1666-74.

41. Kulinski S, Inouye DI. Towards explaining distribution shifts. In: International Conference on Machine Learning. PMLR; 2023. p. 17931-52.

42. Sui Y, Wu Q, Wu J, Cui Q, Li L, Zhou J, et al. Unleashing the power of graph data augmentation on covariate distribution shift. Advances in Neural Information Processing Systems. 2024;36.

43. Zhang M, Marklund H, Dhawan N, Gupta A, Levine S, Finn C. Adaptive risk minimization: Learning to adapt to domain shift. Advances in Neural Information Processing Systems. 2021;34:23664-78.

44. Zhang M, Levine S, Finn C. Memo: Test time robustness via adaptation and augmentation. Advances in Neural Information Processing Systems. 2022;35:38629-42.

45. Donoho DL, et al. High-dimensional data analysis: The curses and blessings of dimensionality. AMS math challenges lecture. 2000;1(2000):32.

46. Cristianini N. An Introduction to Support Vector Machines and other kernel-based learning methods. Cambridge University Press; 2000.

47. Storey JD, Taylor JE, Siegmund D. Strong control, conservative point estimation and simultaneous conservative consistency of false discovery rates: a unified approach. Journal of the Royal Statistical Society Series B: Statistical Methodology. 2004;66(1):187-205.

48. LeCun Y, Bottou L, Bengio Y, Haffner P. Gradient-Based Learning Applied to Document Recognition. Proceedings of the IEEE. 1998;86(11):2278-324. doi:10.1109/5.726791.

49. Viviano JD, Simpson B, Dutil F, Bengio Y, Cohen JP. Saliency is a Possible Red Herring When Diagnosing Poor Generalization. In: International Conference on Learning Representations; 2021. Available from: https://openreview.net/forum?id=c9-WeM-ceB.

50. Veeling BS, Linmans J, Winkens J, Cohen T, Welling M. Rotation Equivariant CNNs for Digital Pathology. 2018 Jun. arXiv:1806.03962.

51. Ehteshami Bejnordi B, Veta M, Johannes van Diest P, van Ginneken B, Karssemeijer N, Litjens G, et al. Diagnostic Assessment of Deep Learning Algorithms for Detection of Lymph Node Metastases in Women With Breast Cancer. JAMA. 2017 12;318(22):2199-210. Available from: https://doi.org/10.1001/jama.2017.14585. arXiv:https://jamanetwork.com/journals/jama/articlepdf/2665774/jama_ehteshami_bejnor doi:10.1001/jama.2017.14585.

52. Pocock J, Graham S, Vu QD, Jahanifar M, Deshpande S, Hadjigeorghiou G, et al. TIAToolbox as an end-to-end library for advanced tissue image analytics. Communications Medicine. 2022 sep;2(1):120. Available from: https://www.nature.com/articles/s43856-022-00186-5. doi:10.1038/s43856-022-00186-5.

53. Irvin J, Rajpurkar P, Ko M, Yu Y, Ciurea-Ilcus S, Chute C, et al. CheXpert: A Large Chest Radiograph Dataset with Uncertainty Labels and Expert Comparison. In: Proceedings of the AAAI Conference on Artificial Intelligence. vol. 33; 2019. p. 590-7. doi:10.1609/aaai.v33i01.3301590.

54. Cohen JP, Hashir M, Brooks R, Bertrand H. On the limits of cross-domain generalization in automated X-ray prediction. In: Medical Imaging with Deep Learning; 2020. Available from: https://arxiv.org/abs/2002.02497.

55. Cohen JP, Viviano JD, Bertin P, Morrison P, Torabian P, Guarrera M, et al. TorchXRayVision: A library of chest X-ray datasets and models. In: Medical Imaging with Deep Learning; 2022. Available from: https://github.com/mlmed/torchxrayvision.

56. Zaki HA, Albaroudi B, Shaban EE, Shaban A, Elgassim M, Almarri ND, et al. Advancement in pleura effusion diagnosis: a systematic review and meta-analysis of point-of-care ultrasound versus radiographic thoracic imaging. The Ultrasound Journal. 2024 Jan;16(1):3. Available from: https://doi.org/10.1186/s13089-023-00356-z. doi:10.1186/s13089-023-00356-z.

57. Cashen K, Petersen TL. Pleural Effusions and Pneumothoraces. Pediatrics in Review. 2017;38:170 181. Available from: https://api.semanticscholar.org/CorpusID:45081803.

58. Mahesh M, Ansari AJ, Mettler Jr FA. Patient exposure from radiologic and nuclear medicine procedures in the United States and worldwide: 2009–2018. Radiology. 2022;307(1):e221263.

59. Shebl E, Paul M. Parapneumonic pleural effusions and empyema thoracis. 2018.

60. Greenwald NF, Miller G, Moen E, Kong A, Kagel A, Dougherty T, et al. Whole-Cell Segmentation of Tissue Images with Human-Level Performance Using Large-Scale Data Annotation and Deep Learning. Nature Biotechnology. 2022;40(4):555-65. doi:10.1038/s41587-021-01094-0.

61. Stringer C, Michaelos M, Pachitariu M. Cellpose: A generalist algorithm for cellular segmentation. 2020 02. doi:10.1101/2020.02.02.931238.

62. Amgad M, Elfandy H, Hussein H, Atteya LA, Elsebaie MAT, Elnasr LSA, et al. Structured crowdsourcing enables convolutional segmentation of histology images. Bioinformatics. 2019;35:3461 3467. Available from: https://api.semanticscholar.org/CorpusID:73424650.

Click here to access/download;Figure;figure_1.pdf

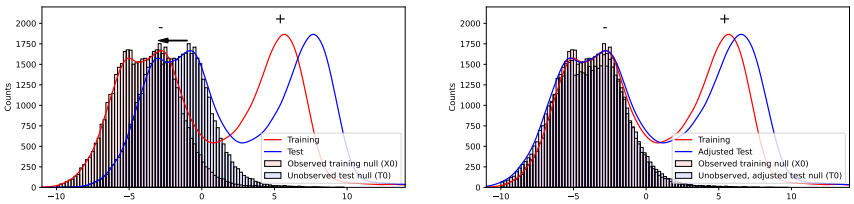


A B

Figure 1: **Sketch of TNA.** (A) The plot illustrates the distribution shift of the discriminative scores of the test (blue) instances from the training (red) instances. The shaded areas illustrate the observed training and unobserved test null distributions, respectively. The black arrow indicates that the test scores should be shifted to leftward. The TNA method aims to approximate the unknown test null distribution (blue shaded area) based on the known true training null distribution (red shaded area). (B) TNA aims to perform a test score adjustment so that the test null aligns to training null distributions better.

Click here to access/download;Figure;figure_2.pdf

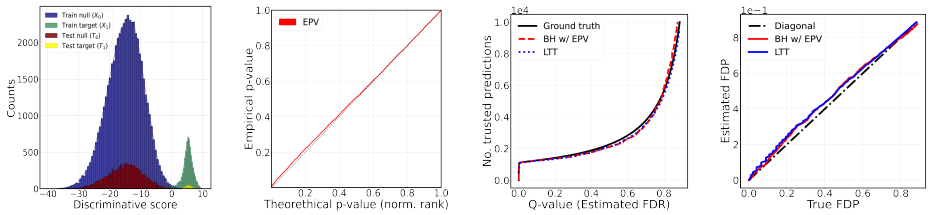


Figure 2: **FDR control in MNIST with EPVs.** (A) The discriminative score distributions. (B) The Q-Q plot of the EPVs of the test samples against the theoretical uniform distribution. (C) The number of the accepted (trusted) positive predictions at various q-values calculated with true labels (black solid line), BH with EPVs (red dashed line), and LTT (blue dotted line). (D) Deviation between the true and the estimated FDR produced by BH with EPV (red line) and LTT (blue dots), respectively.

Click here to access/download;Figure;figure_3.pdf

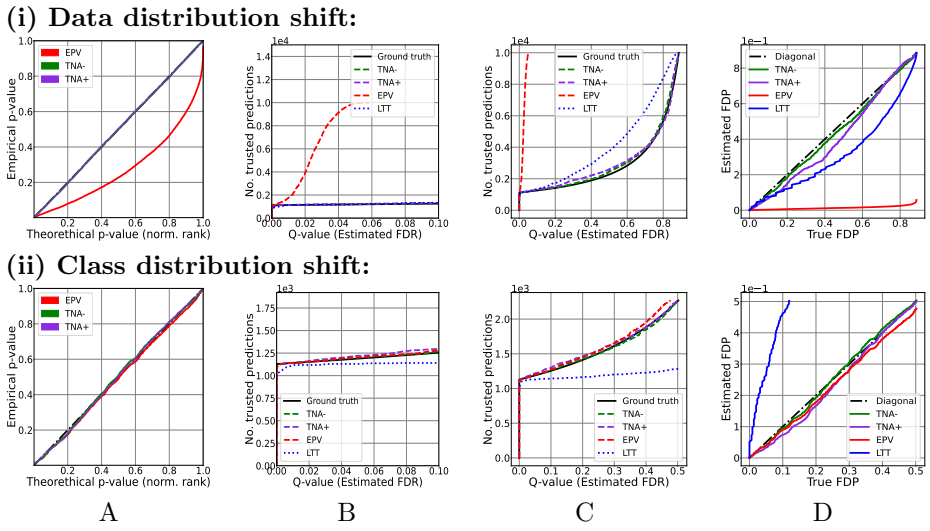


Figure 3: **FDR control under data (top row) and label (bottom row) shift.** (A) Q-Q plot of the EPVs (red dots), the TNA- EPVs (green dots) and the TNA+ EPVs (purple dots) against the theoretical uniform distribution (normalized rank). (B) The number of accepted classifications as a function of the Q-values over a critical range (0-0.1) obtained with (i) ground truth (black line), (ii) BH with EPV (red), (iii) BH with TNA- EPV (green), (iv) BH with TNA- EPV (purple), and (v) LTT (blue). (C) Same as (B) but over the entire q-value range. (D) Deviation of the estimated FDP from the true FDP obtained with (i) BH with EPV (red line), (ii) BH with TNA- EPV (green line), (iii) BH with TNA+ EPV (purple line), and (iv) LTT (blue line).

Click here to access/download;Figure;figure_4.pdf

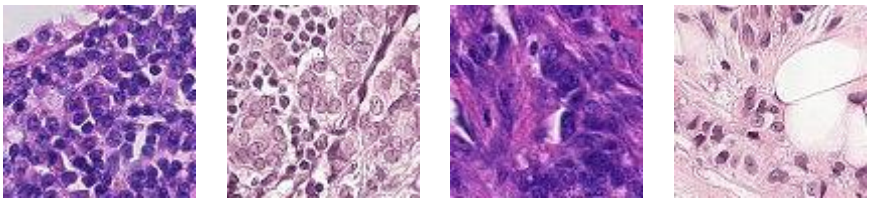

Figure 4: **Instances from PCam dataset.**

Click here to access/download;Figure;figure_5.pdf

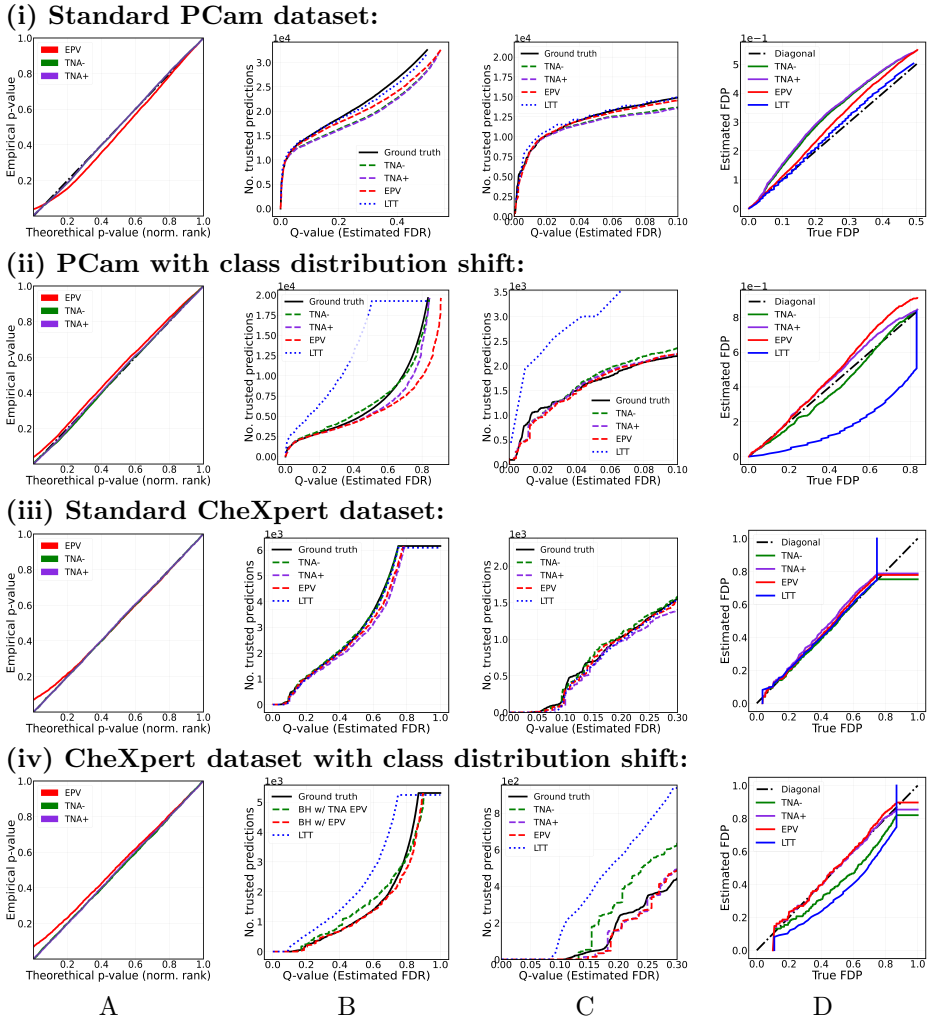


Figure 5: **Performance evaluation of FDR controlling methods.** See the caption of Fig **??** for plot descriptions.

Click here to access/download;Figure;figure_6.pdf

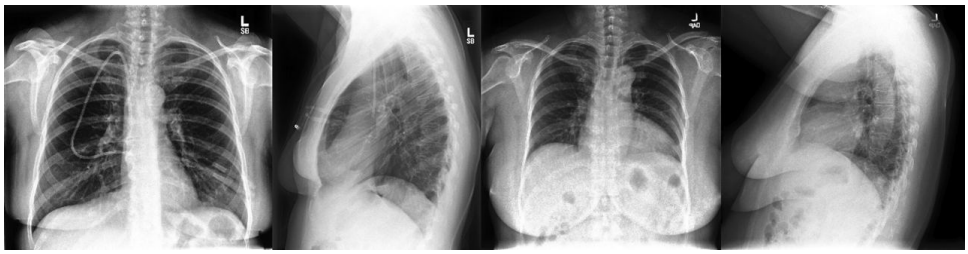

Figure 6: **Instances of CheXpert dataset.**

Click here to access/download;Figure;figure_7.pdf

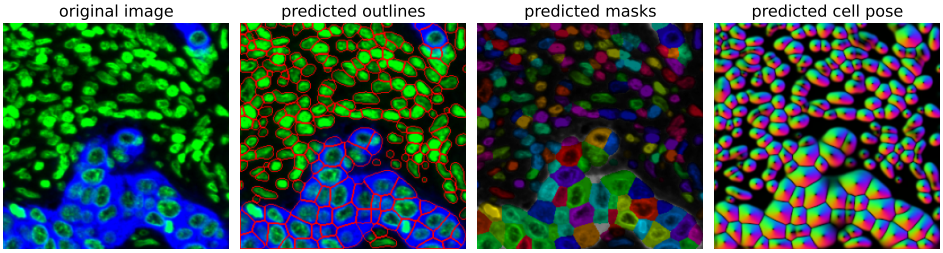


Figure 7: **Instance of TissueNet dataset.** Two-channel original images are being processed by the model, in turn producing masks & flows to correctly segment cells

Click here to access/download;Figure;figure_8.pdf

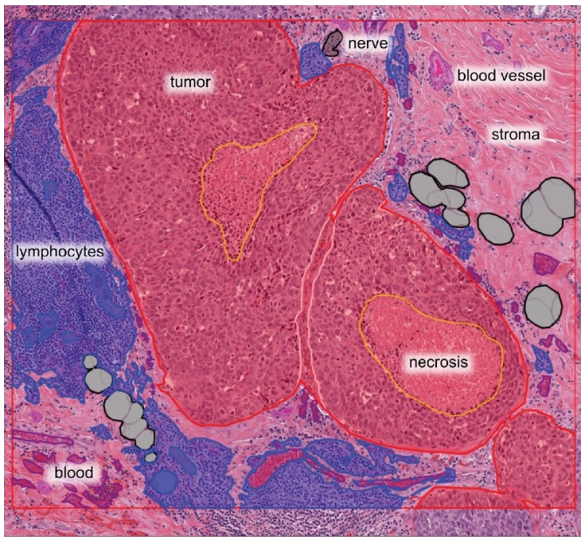


Figure 8: **An instance from BCSS dataset.**

Click here to access/download;Figure;figure_9.pdf

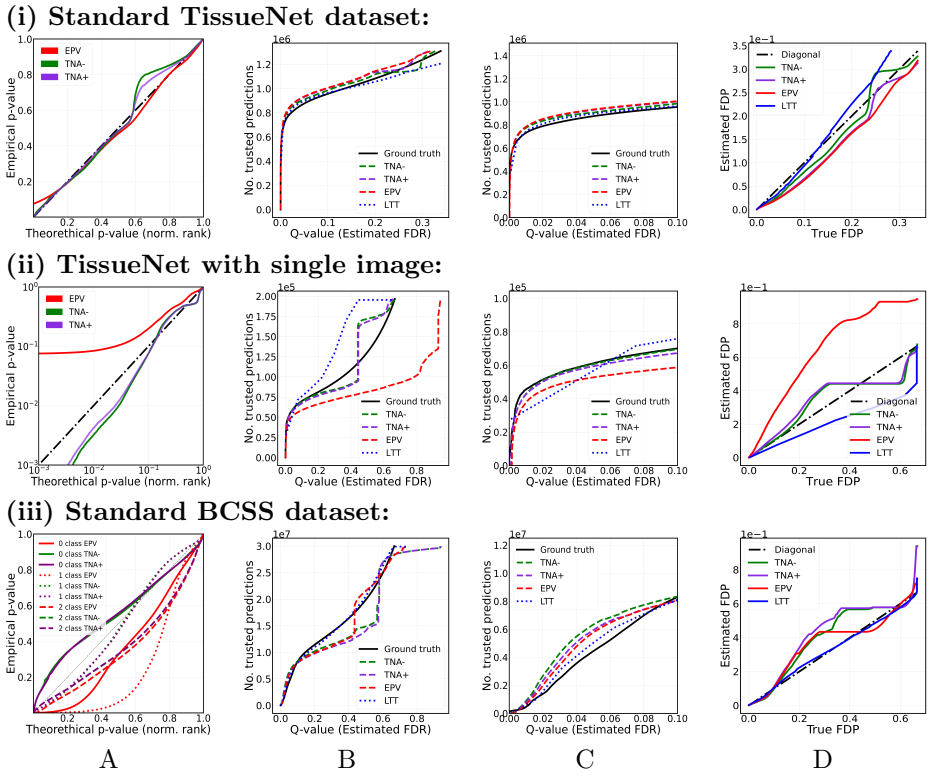


Figure 9: **Performance evaluation of FDR controlling methods.** See the caption of Fig **??** for plot descriptions.

Click here to access/download;Figure;figure_10.pdf

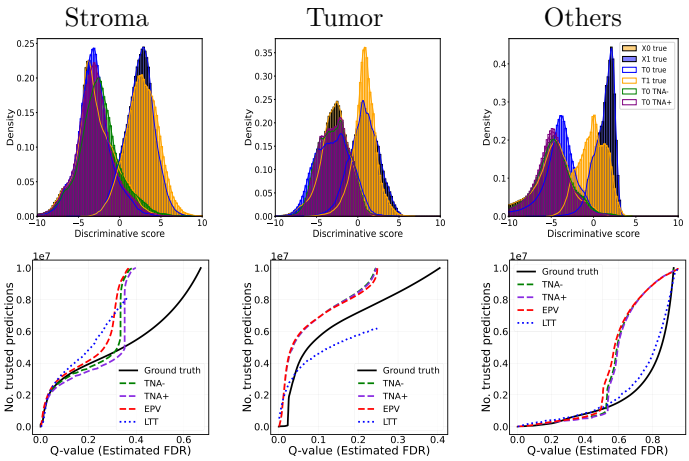


Figure 10: **Classwise score distributions and FDR control in the BCSS dataset.**